

## Topic 12

# Cache vs. hash

資料結構與程式設計  
Data Structure and Programming

12/12/2018

## From $O(\log n)$ to $O(1)$ ?

- ◆ For set and map, they have good complexity for “insert”, “delete” and “find” operations  
→  $O(\log n)$
- ◆ However, in set and map, all the data are sorted --
  - Can output the data in ascending/descending order
  - Can get the list of elements with values in certain range
- ◆ What if we don't care about the order, and just want to have fast “insert”, “delete” and “find” operations?
  - Can we gain something (complexity) back for not sorting the data?


Data Structure and Programming

Prof. Chung-Yang (Ric) Huang

2

## Hash Table

1. Buckets: the table is composed of  $B$  buckets (usually a large number)
2. Each bucket can hold up to  $S$  slots of data (usually a smaller number)
3. Given a data  $d$  with key  $x$ , a hash function  $h(x)$  is used to compute the corresponding bucket number




Data Structure and Programming

Prof. Chung-Yang (Ric) Huang

3

## Hash Table Example

- ◆ Record: (student name, score)
- ◆ Hash table: 26 buckets
- ◆ Hash function = the first character of name



Data Structure and Programming

Prof. Chung-Yang (Ric) Huang

4

## Complexity Analysis

- ◆ Depending on how the  $s$  slots are designed
- ◆ However, the worse case...
  - Insert:  $O(1)$ 
    - Assuming it takes  $O(1)$  to compute  $h(x)$
  - Delete:  $O(s)$  → can they be  $O(\log s)$ ?
  - Find:  $O(s)$  → what's the price to pay?
- ◆ Because  $s$  is usually a smaller number (e.g. 2) → Very efficient

## Hash Table Design Issues

1. Choice of hash function
2. Overflow handling methods
3. Size (number of buckets) of hash table

## Hash Functions

- ◆ Convert key  $x$  to an integer  $b$  as the bucket index ( $0 \leq b \leq B - 1$ )
  - Generally  $O(1)$  complexity
- ◆ Discuss: how's the hash function used in slide #4 (student name, score) example?
  - No good, the first characters of names are usually not evenly distributed among 26 letters
  - Some buckets may become full easily (e.g. 'c'), while some may be empty (e.g. 'x')
- ◆ Ideal: for all possible key values, approximately the same number of keys get mapped into each bucket  
→ Uniform Hash Function

## Hash Function Methods

- ◆ Summation
  - e.g. Adding the ASCII values of some/all the characters together
- ◆ Shift
  - e.g. Keys are pointers; because pointer values are multiplier of 4 (or 8)  
→  $h(x) = (x \gg 2) \dots$
  - e.g.  $x.p$ ,  $x.q$  are two pointers;  
→  $h(x) = (x.p \gg 3) + (x.q \gg 6);$
- ◆ Division
  - e.g. Divide a prime number
- ◆ Others: folding, mid-square, digit analysis, etc
- ◆ Usually: mixed of the above

## Collision and Overflow

- ◆ Collision
    - Two non-identical keys are hashed into the same bucket
    - At most  $(s - 1)$  collisions for a buckets
    - Reduced by better hash function
  - ◆ Overflow
    - When a new key is hashed into a full bucket
- For better hash performance, we should try to produce less collisions and prevent overflow

## Overflow Handling

- ◆ Overflow may still happen when more and more data are stored into the hash
1. Open addressing
    - Find a non-full bucket to insert the new key
    - Linear probing, quadratic probing, rehashing, (pseudo)random probing, etc
  2. Chaining
    - Use linked list, dynamic array, or other kinds of ADT to make the s extendible

## Hash Table Size

- ◆ Hash table size (number of buckets) also affects the occurrence of overflow
    - Too small → Overflow happens
    - Too large → Waste of memory
1. Static hashing
    - Fixed-size hash table
    - Easier to implement; better if the number of possible elements is known in advance
  2. Dynamic hashing
    - Hash table size can grow when necessary
    - When it grows, rehashing is needed

## Hash Classes in STL

- ◆ Static hashing, use set or map for the storage of each bucket
1. hash\_set
  2. hash\_multiset
  3. hash\_map
  4. hash\_multimap
- ◆ However, hash is NOT included as a standard package in newer platforms. You may need to do (For example) ---
    - `#include <hash_set.h>` and/or
    - `g++ -I/usr/include/c++/4.0.0/backward/`

## class hash\_set in STL

- ◆ hash\_set<Key[, HashFunc, EqualKey, Alloc]>
  - class Key: element type
  - class HashFunc: hash function (optional; default = hash<Key>)
  - class EqualKey: equality checking for class Key (optional; default = equal\_to<Key>)
  - class Alloc: used for internal memory management (optional; default = alloc)

## Member Functions of class hash\_set

1. iterator begin() const;  
iterator end() const;
2. size\_type bucket\_count() const;
3. void resize(size\_type n);
4. pair<iterator, bool> insert(const value\_type& x);  
void insert(InputIterator f, InputIterator l);
5. void erase(iterator pos);  
size\_type erase(const key\_type& k);  
void erase(iterator first, iterator last);
6. iterator find(const key\_type& k) const;
7. size\_type count(const key\_type& k) const;

## Other Hash Classes in STL

- ◆ class hash\_multiset in STL
  - Similar to hash\_set, but allow elements with identical keys
- ◆ class hash\_map
  - hash\_map<Key, Data[, HashFunc, EqualKey, Alloc]>
- ◆ class hash\_multimap
  - Similar to hash\_map, but allow elements with identical keys

## Hash Implementation Example (myHashSet.h)

```
template <class Data>
class HashSet
{
public:
    class iterator
    {
        friend class HashSet<Data>;
    };
private:
    size_t          _numBuckets;
    // Each _buckets[i] is a vector<Data>
    vector<Data>*   _buckets;
};
```

Must overload:

- operator () → as the hash function (i.e. d() % \_numBuckets to return bucket index)
- operator == → to compare if the data are equivalent (or have the same key)

## Supported functions in class HashSet

```

iterator begin() const; // Point to the first valid data
iterator end() const; // Past the end

bool empty() const; // return true if no valid data
size_t size() const; // number of valid data
vector<Data>& operator [] (size_t i) { return _buckets[i]; }

const vector<Data>& operator [] (size_t i) const;
void init(size_t b); // initialize Hash with _numBuckets = b
void reset();

```

## Supported functions in class HashSet

```

// check if d is in the hash...
// if yes, return true; else return false;
bool check(const Data& d) const;

// query if d is in the hash...
// if yes, replace d with the data in the hash and return true;
// else return false;
bool query(Data& d) const;

// update the entry in hash that is equal to d (i.e. == return true)
// if found, update that entry with d and return true;
// else insert d into hash as a new entry and return false;
bool update(const Data& d);

// return true if inserted successfully (i.e. d is not in the hash)
// return false if d is already in the hash ==> will not insert
bool insert(const Data& d);

```

## Another Hash Implementation Example (myHashMap.h)

```

template <class HashKey, class HashData>
class HashMap
{
    typedef pair<HashKey, HashData> HashNode;
public:
    class iterator
    {
        friend class HashMap<HashKey, HashData>;
    };
private:
    size_t          _numBuckets;
    vector<HashNode>* _buckets;
};

```

## Another Hash Implementation Example (myHashMap.h)

class HashMap<HashKey, HashData>

vector<HashNode>\* \_buckets

\_numBuckets

vector<HashNode>

All the HashNodes have the same bucketNum(), but distinct HashKeys

HashNode  
≡ pair<HashKey, HashData> (k, d)

Overload operator () for HashKey as hash function

Usually a prime number (Why?)

size\_t bucketNum(const HashKey& k) const { return (k()) % \_numBuckets; }

## Class HashKey

- ◆ To use Hash ADT, you should define your own HashKey class.
- It should at least overload the "()" and "==" operators.

```
class HashKey // Of course, name your own HashKey
{
public:
    HashKey(); // define your own constructor
    size_t operator() () const; // acted as "hash
                                // function"

    bool operator == (const HashKey& k);
                                // to compare the HashKey

private:
    // Define your own data members
};
```

## Example of using class Hash

- ◆ Locating an address
- ◆ `typedef string Address;`
- ◆ `typedef pair<float, float> Location;`
- ◆ `class AddressKey {`
- ◆ `public:`
- ◆ `size_t operator() () const { ...; }`
- ◆ `bool operator == (const Address& a) {`
- ◆ `return (_addr == a._addr); }`
- ◆ `private:`
- ◆ `Address _addr;`
- ◆ `};`

## class Hash::iterator

```
class iterator
{
    friend class HashMap<HashKey, HashData>;
private:
    // Define your own data members!!
};
◆ Purpose: to go through the "valid" HashNodes in the Hash
◆ To use:
    HashMap<HashKey, HashData> hh;
    ... // insert data
    HashMap<HashKey, HashData>::iterator hi = hh.begin();
    for (; hi != hh.end(); ++hi)
        cout << (*hi).first << " → " << (*hi).second
               << endl;
```

## Supported functions in class HashMap

```
iterator begin() const; // Point to the first valid data
iterator end() const; // Pass the end
bool empty() const; // return true if no valid data
size_t size() const; // number of valid data
vector<HashNode>& operator [] (size_t i) { return _buckets[i]; }
const vector<HashNode>& operator [] (size_t i) const;
void init(size_t b); // initialize Hash with _numBuckets = b
void reset();
bool check(const HashKey& k, HashData& n);
bool insert(const HashKey& k, const HashData& d);
bool replaceInsert(const HashKey& k, const HashData& d);
void forceInsert(const HashKey& k, HashData d);
```

## Cache in Programming

- ◆ Structurally similar to hash, however
  - Usually smaller number of buckets
  - Each bucket contains exactly 1 element
  - When collision happens, the old data is overwritten by the new one
  - Easier to implement than hash
- ◆ Usually used as computational cache
  - (input parameters) → computational results
- ◆ There is no STL implement yet

## Cache Implementation in util/myHashMap.h

```
template <class CacheKey, class CacheData>
class Cache
{
    typedef pair<CacheKey, CacheData> CacheNode;
public:
    // No need to implement class iterator (why?)
    void init(size_t s);
    void reset();
    size_t size() const;
    CacheNode& operator [] (size_t i);
    const CacheNode& operator [] (size_t i) const;
    bool read(const CacheKey& k, CacheData& d) const;
    void write(const CacheKey& k, const CacheData& d);
private:
    size_t      _size;
    CacheNode*  _cache; // new CacheNode[_size]
};
```

## Example of using class Cache (BDD project)

- ◆ Computed table
$$\text{ITE}(F, G, H) = F * G + \bar{F} * H$$

F, G, H: BddNode (contains a size\_t data)

  - Requires expensive recursive calls to compute ITE() functions ( $O(|H|*|G|*|H|)$ )
  - The computed table (cache) is to record the result (as CacheData) with respect to the ITE parameters (as CacheKey)
  - So next time when the same ITE(F, G, H) computation is required, we can immediately look up the cached result

## FYI: Google Hash

<http://code.google.com/p/google-sparsehash/>

- ◆ Two hashtable implementations:
  1. Sparse hash: designed to be space efficient
  2. Dense hash: designed to be time efficient
- ◆ Interface very similar to SGI's (STL) hash\_map, hash\_set, etc. But is claimed to be much more efficient.
  - sparse\_hash\_map, sparse\_hash\_set, dense\_hash\_map, dense\_hash\_set
  - e.g. sparse\_hash\_map<Key, Data, HashFcn, EqualKey, Alloc>